

# Learning Personalized End-To-End Goal-Oriented Dialog

**Liangchen Luo**

Peking University, Beijing, China  
luolc@pku.edu.cn

**Wenhao Huang**

Microsoft Research, Beijing, China  
wenhhu@microsoft.com

## Abstract

This document contains the instructions for preparing a camera-ready manuscript for the proceedings of ACL 2018. The document itself conforms to its own specifications, and is therefore an example of what your manuscript should look like. These instructions should be used for both papers submitted for review and for final versions of accepted papers. Authors are asked to conform to all the directions reported in this document.

## 1 Introduction

As conversational agents develop rapidly, there has been growing research interest in training dialog systems with end-to-end models (Vinyals and Le, 2015; Sordoni et al., 2015; Bordes et al., 2016; Sukhbaatar et al., 2015), which are directly trained on past dialogs without assumption on the domain or dialog state structure. One major issue for these content-based systems is their propensity to select the response only relying on the content of conversation context, overlooking the differences among interlocutors. Therefore, existing models are not aware or capable of adapting to different kinds of users, as illustrated in Table 1. In the first and second case, the user utterances are the same, while one is from a young female and the other from an elderly male. Content-based models cannot adjust speech style flexibly according to who it is interacting with. The third case shows that in a recommendation scenario, users may prefer a specific candidate owing to their personal preferences.

When conversing, humans often adapt to their interlocutor to facilitate understanding, and thus improve conversational efficiency and satisfaction (Serban et al., 2015). Attaining smart and human-level performance with dialog agents may well re-

user info	gender: female, age: young
message	Good morning.
content-based	Hello, what can I help you with today?
personalized	Hey girl! How is it going?
user info	gender: male, age: elderly
message	Good morning.
content-based	Hello, what can I help you with today?
personalized	Greetings sir. What may I assist you with today?
user info	dietary: vegetable
message	Can you make a restaurant reservation in a cheap price range in Rome for six people with Spanish food?
candidates	- rome_cheap_spanish_1 R.type non-veg - rome_cheap_spanish_2 R.type veg
content-based	How about this: rome_cheap_spanish_1?
personalized	How about this: rome_cheap_spanish_2?

Table 1: Three example cases chosen from personalized bAbI dialogs dataset to show the necessity of personalization in dialog systems. Content-based and personalized responses in the examples are generated by a memory network model and a personalized model described in this paper, respectively.

quire personalization. For example, a restaurant reservation agent should be aware of the user profile (such as dietary preference and favorite food), and then make recommendations accordingly. The language style used by agents may be adjusted based on certain characteristics of the interlocutor (age, gender, etc.) as well. Such capabilities could increase the effectiveness and naturalness of generated dialogs (Lucas et al., 2009; Su et al., 2013).

There has been initial studies on endowing bots with a specific personality to improve the response consistency, but without being aware of the user profile (Li et al., 2016; Qian et al., 2017). These works do not focus on enhancing effectiveness by using the information of interlocutor, as we do in this paper. We address the challenge of profile

modeling and how to endow goal-oriented dialog systems with the user profile. For such purposes, our focus is on explicit profile information such as attributes represented by key-value pairs, not latent variables.

Fortunately, neural models of next response retrieval such as end-to-end memory network (Sukhbaatar et al., 2015; Bordes et al., 2016) provide an easy mechanism for incorporating profiles as embeddings. We hence explore a PrO-file EMbedding model (POEM MODEL), within an end-to-end memory network framework. We use the MEMN2N architecture of Sukhbaatar et al. (2015), integrating a profile vector representation into the query and ranking part of the MEMN2N model. Moreover, we also attempt a PersonAL-ized MemorY model (PALMY MODEL), which is no need of explicit profile attributes. The PALMY MODEL follows the vein of MEMN2N as well, with an additional global memory storing conversation context from other users with similar profile. The two models and their combination are trained on human-agent conversation data and used at test time to retrieve personalized responses. Our experiments on goal-oriented dialog corpus of personalized bAbI dialog datasets show that leveraging persona information can significantly improve the performance of memory network. POEM MODEL, PALMY MODEL and their combination all outperform the current state-of-the-art method, where POEM MODEL takes the best with over 8% improvement to the best baseline in per-response accuracy score.

## 2 Related Work

Considerable research efforts have been devoted so far to make conversational agents smarter. One of the approaches is integrating user/speaker model into dialog systems. In this way, an agent can learn abilities such as: choosing the best response according to the specific user preference (Joshi et al., 2017); and using appropriate speech style based on profile of who it is interacting with (Herzig et al., 2017).

A first attempt to model persona is from Li et al. (2016), who propose to assign specific personality and conversation style to agents based on learned persona embeddings. Recently, Luan et al. (2017) describe an interesting approach in the use of multi-task learning with personalized text data. Other researchers attempt to introduce personal-

ized information to dialog systems by transfer learning (Yang et al., 2017; Zhang et al., 2017).

Since there is usually no explicit personalized information in conversation context, existing models (Qian et al., 2017; Herzig et al., 2017) often require extra procedures to incorporate personalization in training, such as intermediate supervision and pre-training of user profile. The present work, by contrast, is totally end-to-end and data-driven. Instead of assigning a personality to agents (Li et al., 2016; Luan et al., 2017; Qian et al., 2017), our model pays more attention on the user persona and aim to make agents more adaptive to different kinds of interlocutors.

It is well accepted that conversation agents include goal-oriented dialog systems and non goal-oriented (*chit-chat*) bots. As most of the previous works (Li et al., 2016; Liu et al., 2018; Qian et al., 2017) focus on *chit-chat*, the combination of personalization and goal-oriented dialog remains to be explored. There are no open datasets which allow researchers to train goal-oriented dialog with personalized information, as argued by Joshi et al. (2017), until they released one. Our work is in the vein of the memory network models for goal-oriented dialog from Sukhbaatar et al. (2015) and Bordes et al. (2016), enriching these models by incorporating profile embeddings and using conversation context from users with similar attributes as global memory.

## 3 Model

## 4 Dataset

## 5 Experiments

## 6 Analysis

## 7 Conclusion

## 8 Credits

This document has been adapted from the instructions for earlier ACL and NAACL proceedings, including those for ACL 2017 by Dan Gildea and Min-Yen Kan, NAACL 2017 by Margaret Mitchell, ACL 2012 by Maggie Li and Michael White, those from ACL 2010 by Jing-Shing Chang and Philipp Koehn, those for ACL 2008 by Johanna D. Moore, Simone Teufel, James Allan, and Sadaoki Furui, those for ACL 2005 by Hwee Tou Ng and Kemal Oflazer, those for ACL 2002 by Eugene Charniak and Dekang Lin, and earlier ACL

and EACL formats. Those versions were written by several people, including John Chen, Henry S. Thompson and Donald Walker. Additional elements were taken from the formatting instructions of the *International Joint Conference on Artificial Intelligence* and the *Conference on Computer Vision and Pattern Recognition*.

## 9 Introduction

The following instructions are directed to authors of papers submitted to ACL 2018 or accepted for publication in its proceedings. All authors are required to adhere to these specifications. Authors are required to provide a Portable Document Format (PDF) version of their papers. **The proceedings are designed for printing on A4 paper.**

## 10 General Instructions

Manuscripts must be in two-column format. Exceptions to the two-column format include the title, authors' names and complete addresses, which must be centered at the top of the first page, and any full-width figures or tables (see the guidelines in Subsection 10.6). **Type single-spaced.** Start all pages directly under the top margin. See the guidelines later regarding formatting the first page. The manuscript should be printed single-sided and its length should not exceed the maximum page limit described in Section 12. Pages are numbered for initial submission. However, **do not number the pages in the camera-ready version.**

By uncommenting `\aclfinalcopy` at the top of this document, it will compile to produce an example of the camera-ready formatting; by leaving it commented out, the document will be anonymized for initial submission. When you first create your submission on softconf, please fill in your submitted paper ID where `***` appears in the `\def\aclpaperid{***}` definition at the top.

The review process is double-blind, so do not include any author information (names, addresses) when submitting a paper for review. However, you should maintain space for names and addresses so that they will fit in the final (accepted) version. The ACL 2018  $\LaTeX$  style will create a titlebox space of 6.35 cm for you when `\aclfinalcopy` is commented out.

### 10.1 The Ruler

The ACL 2018 style defines a printed ruler which should be present in the version submitted for review. The ruler is provided in order that reviewers may comment on particular lines in the paper without circumlocution. If you are preparing a document without the provided style files, please arrange for an equivalent ruler to appear on the final output pages. The presence or absence of the ruler should not change the appearance of any other content on the page. The camera ready copy should not contain a ruler. ( $\LaTeX$  users may uncomment the `\aclfinalcopy` command in the document preamble.)

**Reviewers:** note that the ruler measurements do not align well with lines in the paper – this turns out to be very difficult to do well when the paper contains many figures and equations, and, when done, looks ugly. In most cases one would expect that the approximate location will be adequate, although you can also use fractional references (e.g., the first paragraph on this page ends at mark 117.5).

### 10.2 Electronically-available resources

ACL provides this description in  $\LaTeX$ 2e (`acl2018.tex`) and PDF format (`acl2018.pdf`), along with the  $\LaTeX$ 2e style file used to format it (`acl2018.sty`) and an ACL bibliography style (`aclnatbib.bst`) and example bibliography (`acl2018.bib`). These files are all available at <http://acl2018.org/downloads/acl18-latex.zip>.

A Microsoft Word template file (`acl18-word.docx`) and example submission pdf (`acl18-word.pdf`) is available at <http://acl2018.org/downloads/acl18-word.zip>. We strongly recommend the use of these style files, which have been appropriately tailored for the ACL 2018 proceedings.

### 10.3 Format of Electronic Manuscript

For the production of the electronic manuscript you must use Adobe's Portable Document Format (PDF). PDF files are usually produced from  $\LaTeX$  using the `pdflatex` command. If your version of  $\LaTeX$  produces Postscript files, you can convert these into PDF using `ps2pdf` or `dvipdf`. On Windows, you can also use Adobe Distiller to generate PDF.

Please make sure that your PDF file includes all the necessary fonts (especially tree diagrams, symbols, and fonts with Asian characters). When you print or create the PDF file, there is usually an option in your printer setup to include none, all or just non-standard fonts. Please make sure that you select the option of including ALL the fonts. **Before sending it, test your PDF by printing it from a computer different from the one where it was created.** Moreover, some word processors may generate very large PDF files, where each page is rendered as an image. Such images may reproduce poorly. In this case, try alternative ways to obtain the PDF. One way on some systems is to install a driver for a postscript printer, send your document to the printer specifying “Output to a file”, then convert the file to PDF.

It is of utmost importance to specify the **A4 format** (21 cm x 29.7 cm) when formatting the paper. When working with dvips, for instance, one should specify `-t a4`. Or using the command `\special{papersize=210mm,297mm}` in the latex preamble (directly below the `\usepackage` commands). Then using `dvipdf` and/or `pdflatex` which would make it easier for some.

Print-outs of the PDF file on A4 paper should be identical to the hardcopy version. If you cannot meet the above requirements about the production of your electronic submission, please contact the publication chairs as soon as possible.

## 10.4 Layout

Format manuscripts two columns to a page, in the manner these instructions are formatted. The exact dimensions for a page on A4 paper are:

- Left and right margins: 2.5 cm
- Top margin: 2.5 cm
- Bottom margin: 2.5 cm
- Column width: 7.7 cm
- Column height: 24.7 cm
- Gap between columns: 0.6 cm

Papers should not be submitted on any other paper size. If you cannot meet the above requirements about the production of your electronic submission, please contact the publication chairs above as soon as possible.

Type of Text	Font Size	Style
paper title	15 pt	bold
author names	12 pt	bold
author affiliation	12 pt	
the word “Abstract”	12 pt	bold
section titles	12 pt	bold
subsection titles	11 pt	bold
document text	11 pt	
captions	11 pt	
abstract text	11 pt	
bibliography	10 pt	
footnotes	9 pt	

Table 2: Font guide.

## 10.5 Fonts

For reasons of uniformity, Adobe’s **Times Roman** font should be used. In  $\text{\LaTeX}2\text{e}$  this is accomplished by putting

```
\usepackage{times}
\usepackage{latexsym}
```

in the preamble. If Times Roman is unavailable, use **Computer Modern Roman** ( $\text{\LaTeX}2\text{e}$ ’s default). Note that the latter is about 10% less dense than Adobe’s Times Roman font.

## 10.6 The First Page

Center the title, author name(s), and affiliation(s) across both columns (or, for the initial submission, **Anonymous ACL submission** for names and affiliations). Do not use footnotes for affiliations. Include the paper ID number assigned during the submission process in the header. Use the two-column format only when you begin the abstract.

**Title:** Place the title centered at the top of the first page, in a 15-point bold font. (For a complete guide to font sizes and styles, see Table 2) Long titles should be typed on two lines without a blank line intervening. Approximately, put the title at 2.5 cm from the top of the page, followed by a blank line, then the author name(s), and the affiliation(s) on the following line. Do not use only initials for given names (middle initials are allowed). Do not format surnames in all capitals (*e.g.*, use “Mitchell” not “MITCHELL”). Do not format title and section headings in all capitals as well except for proper names (such as “BLEU”) that are conventionally in all capitals. The affiliation should contain the author’s complete address, and if pos-



Command	Output	Command	Output
<code>{\ "a}</code>	ä	<code>{\c c}</code>	ç
<code>{\ ^e}</code>	ê	<code>{\u g}</code>	ğ
<code>{\ `i}</code>	ì	<code>{\l}</code>	ł
<code>{\ .I}</code>	İ	<code>{\~n}</code>	ñ
<code>{\o}</code>	ø	<code>{\H o}</code>	ö
<code>{\ 'u}</code>	ú	<code>{\v r}</code>	ř
<code>{\aa}</code>	å	<code>{\ss}</code>	ß

Table 3: Example commands for accented characters, to be used in, *e.g.*, BibTeX names.

sible, an electronic mail address. Start the body of the first page 7.5 cm from the top of the page.

The title, author names and addresses should be completely identical to those entered to the electronic paper submission website in order to maintain the consistency of author information among all publications of the conference. If they are different, the publication chairs may resolve the difference without consulting with you; so it is in your own interest to double-check that the information is consistent.

**Abstract:** Type the abstract at the beginning of the first column. The width of the abstract text should be smaller than the width of the columns for the text in the body of the paper by about 0.6 cm on each side. Center the word **Abstract** above the body of the abstract using the font size and style shown in Table 2. The abstract should be a concise summary of the general thesis and conclusions of the paper. It should be no longer than 200 words. The font size of the abstract text should be as shown in Table 2.

**Text:** Begin typing the main body of the text immediately after the abstract, observing the two-column format as shown in the present document. Do not include page numbers in the final version.

**Indent:** Indent when starting a new paragraph, about 0.4 cm.

## 10.7 Sections

**Headings:** Type and label section and subsection headings in the style shown on the present document. Use numbered sections (Arabic numerals) in order to facilitate cross references. Number subsections with the section number and the subsection number separated by a dot, in Arabic numerals. Do not number subsubsections (*i.e.*, use `\subsubsection*` instead of `\subsubsection`).

**Citations:** Citations within the text appear in parentheses as (Gusfield, 1997) or, if the author’s

name appears in the text itself, as Gusfield (1997). Using the provided L<sup>A</sup>T<sub>E</sub>X style, the former is accomplished using `\cite` and the latter with `\shortcite` or `\newcite`. Collapse multiple citations as in (Gusfield, 1997; Aho and Ullman, 1972); this is accomplished with the provided style using commas within the `\cite` command, *e.g.*, `\cite{Gusfield:97,Aho:72}`. Append lowercase letters to the year in cases of ambiguities. Treat double authors as in (Aho and Ullman, 1972), but write as in (Chandra et al., 1981) when more than two authors are involved.

Also refrain from using full citations as sentence constituents. We suggest that instead of

“(Gusfield, 1997) showed that ...”

you use

“Gusfield (1997) showed that ...”

If you are using the provided L<sup>A</sup>T<sub>E</sub>X and BibTeX style files, you can use the command `\citete` (cite in text) to get “author (year)” citations.

You can use the command `\citealp` (alternative cite without parentheses) to get “author year” citations (which is useful for using citations within parentheses, as in Gusfield, 1997).

If the BibTeX file contains DOI fields, the paper title in the references section will appear as a hyperlink to the DOI, using the `hyperref` L<sup>A</sup>T<sub>E</sub>X package. To disable the `hyperref` package, load the style file with the `nohyperref` option: `\usepackage[nohyperref]{acl2018}`.

**Compilation Issues:** Some of you might encounter the following error during compilation:

“`\pdfendlink ended up in different nesting level than \pdfstartlink.`”

This happens when `pdflatex` is used and a citation splits across a page boundary. To fix this, disable the `hyperref` package (see above), recompile and see the problematic citation. Next rewrite that sentence containing the citation. (See, *e.g.*, <http://tug.org/errors.html>)

**Digital Object Identifiers:** As part of our work to make ACL materials more widely used and cited outside of our discipline, ACL has registered as a CrossRef member, as a registrant of Digital Object Identifiers (DOIs), the standard for registering permanent URNs for referencing scholarly materials. We are requiring all camera-ready references to contain the appropriate DOIs (or as a second resort, the hyperlinked ACL Anthology

output	natbib	previous ACL style files
(Gusfield, 1997)	\citep	\cite
Gusfield (1997)	\citet	\newcite
(1997)	\citeyearpar	\shortcite

Table 4: Citation commands supported by the style file. The citation style is based on the natbib package and supports all natbib citation commands. It also supports commands defined in previous ACL style files for compatibility.

Identifier) to all cited works. Thus, please ensure that you use Bib<sub>T</sub> records that contain DOI or URLs for any of the ACL materials that you reference. Appropriate records should be found for most materials in the current ACL Anthology at <http://aclanthology.info/>.

As examples, we cite (Goodman et al., 2016) to show you how papers with a DOI will appear in the bibliography. We cite (Harper, 2014) to show how papers without a DOI but with an ACL Anthology Identifier will appear in the bibliography.

As reviewing will be double-blind, the submitted version of the papers should not include the authors’ names and affiliations. Furthermore, self-references that reveal the author’s identity, e.g.,

“We previously showed (Gusfield, 1997) ...”

should be avoided. Instead, use citations such as

“Gusfield (1997) previously showed ... ”

**Please do not use anonymous citations** and do not include acknowledgments when submitting your papers. Papers that do not conform to these requirements may be rejected without review.

**References:** Gather the full set of references together under the heading **References**; place the section before any Appendices, unless they contain references. Arrange the references alphabetically by first author, rather than by order of occurrence in the text. Provide as complete a citation as possible, using a consistent format, such as the one for *Computational Linguistics* or the one in the *Publication Manual of the American Psychological Association* (American Psychological Association, 1983). Use of full names for authors rather than initials is preferred. A list of abbreviations for common computer science journals can be found in the *ACM Computing Reviews* (Association for Computing Machinery, 1983).

The L<sub>A</sub>T<sub>E</sub>X and Bib<sub>T</sub> style files provided roughly fit the American Psychological Associa-

tion format, allowing regular citations, short citations and multiple citations as described above.

**Appendices:** Appendices, if any, directly follow the text and the references (but see above). Letter them in sequence and provide an informative title: **Appendix A. Title of Appendix**.

## 10.8 Footnotes

**Footnotes:** Put footnotes at the bottom of the page and use the footnote font size shown in Table 2. They may be numbered or referred to by asterisks or other symbols.<sup>1</sup> Footnotes should be separated from the text by a line.<sup>2</sup>

## 10.9 Figures and Tables

**Placement:** Place figures and tables in the paper near where they are first discussed, as close as possible to the top of their respective column.

**Captions:** Provide a caption for every illustration; number each one sequentially in the form: “Figure 1: Caption of the Figure.” “Table 1: Caption of the Table.” Type the captions of the figures and tables below the body, using the caption font size shown in Table 2.

## 10.10 Equation

An example equation is shown below:

$$A = \pi r^2 \quad (1)$$

The numbering (if any) and alignment of the equations will be done automatically (using `align` or `equation`).

## 10.11 Accessibility

In an effort to accommodate the color-blind (as well as those printing to paper), grayscale readability for all accepted papers will be encouraged. Color is not forbidden, but authors should ensure that tables and figures do not rely solely on color to convey critical distinctions. A simple criterion:

<sup>1</sup>This is how a footnote should appear.

<sup>2</sup>Note the line separating the footnotes from the text.

All curves and points in your figures should be clearly distinguishable without color.

## 11 Translation of non-English Terms

It is also advised to supplement non-English characters and terms with appropriate transliterations and/or translations since not all readers understand all such characters and terms. Inline transliteration or translation can be represented in the order of: original-form transliteration “translation”.

## 12 Length of Submission

The ACL 2018 main conference accepts submissions of long papers and short papers. Long papers may consist of up to eight (8) pages of content plus unlimited pages for references. Upon acceptance, final versions of long papers will be given one additional page – up to nine (9) pages of content plus unlimited pages for references – so that reviewers’ comments can be taken into account. Short papers may consist of up to four (4) pages of content, plus unlimited pages for references. Upon acceptance, short papers will be given five (5) pages in the proceedings and unlimited pages for references.

For both long and short papers, all illustrations and tables that are part of the main text must be accommodated within these page limits, observing the formatting instructions given in the present document. Supplementary material in the form of appendices does not count towards the page limit.

However, note that supplementary material should be supplementary (rather than central) to the paper, and that reviewers may ignore supplementary material when reviewing the paper (see Appendix A). Papers that do not conform to the specified length and formatting requirements are subject to be rejected without review.

Workshop chairs may have different rules for allowed length and whether supplemental material is welcome. As always, the respective call for papers is the authoritative source.

## Acknowledgments

The acknowledgments should go immediately before the references. Do not number the acknowledgments section (*i.e.*, use \section\* instead of \section). Do not include this section when submitting your paper for review.

## References

- Alfred V. Aho and Jeffrey D. Ullman. 1972. *The Theory of Parsing, Translation and Compiling*, volume 1. Prentice-Hall, Englewood Cliffs, NJ.
- American Psychological Association. 1983. *Publications Manual*. American Psychological Association, Washington, DC.
- Association for Computing Machinery. 1983. *Computing Reviews* 24(11):503–512.
- Antoine Bordes, Y-Lan Boureau, and Jason Weston. 2016. [Learning end-to-end goal-oriented dialog](#). In *the International Conference on Learning Representations (ICLR)*. <https://openreview.net/forum?id=S1Bb3D5gg>.
- Ashok K. Chandra, Dexter C. Kozen, and Larry J. Stockmeyer. 1981. [Alternation](#). *Journal of the Association for Computing Machinery* 28(1):114–133. <https://doi.org/10.1145/322234.322243>.
- James Goodman, Andreas Vlachos, and Jason Naradowsky. 2016. [Noise reduction and targeted exploration in imitation learning for abstract meaning representation parsing](#). In *Proceedings of the 54th Annual Meeting of the Association for Computational Linguistics (Volume 1: Long Papers)*. Association for Computational Linguistics, pages 1–11. <https://doi.org/10.18653/v1/P16-1001>.
- Dan Gusfield. 1997. *Algorithms on Strings, Trees and Sequences*. Cambridge University Press, Cambridge, UK.
- Mary Harper. 2014. [Learning from 26 languages: Program management and science in the babel program](#). In *Proceedings of COLING 2014, the 25th International Conference on Computational Linguistics: Technical Papers*. Dublin City University and Association for Computational Linguistics, page 1. <http://aclweb.org/anthology/C14-1001>.
- Jonathan Herzig, Michal Shmueli-Scheuer, Tommy Sandbank, and David Konopnicki. 2017. [Neural response generation for customer service based on personality traits](#). In *Proceedings of the 10th International Conference on Natural Language Generation*. Association for Computational Linguistics, Santiago de Compostela, Spain, pages 252–256. <http://www.aclweb.org/anthology/W17-3541>.
- Chaitanya K. Joshi, Fei Mi, and Boi Faltings. 2017. [Personalization in goal-oriented dialog](#). *CoRR* abs/1706.07503. <http://arxiv.org/abs/1706.07503>.
- Jiwei Li, Michel Galley, Chris Brockett, Georgios Spithourakis, Jianfeng Gao, and Bill Dolan. 2016. [A persona-based neural conversation model](#). In *Proceedings of the 54th Annual Meeting of the Association for Computational Linguistics (Volume 1: Long Papers)*. Association for Computational Linguistics, Berlin, Germany, pages 994–1003. <http://www.aclweb.org/anthology/P16-1094>.

- B. Liu, Z. Xu, C. Sun, B. Wang, X. Wang, D. F. Wong, and M. Zhang. 2018. [Content-oriented user modeling for personalized response ranking in chatbots](#). *IEEE/ACM Transactions on Audio, Speech, and Language Processing* 26(1):122–133. <https://doi.org/10.1109/TASLP.2017.2763243>.
- Yi Luan, Chris Brockett, Bill Dolan, Jianfeng Gao, and Michel Galley. 2017. [Multi-task learning for speaker-role adaptation in neural conversation models](#). *CoRR* abs/1710.07388. <http://arxiv.org/abs/1710.07388>.
- JM Lucas, F Fernández, J Salazar, J Ferreiros, and R San Segundo. 2009. Managing speaker identity and user profiles in a spoken dialogue system. *Procesamiento del lenguaje natural* (43).
- Qiao Qian, Minlie Huang, Haizhou Zhao, Jingfang Xu, and Xiaoyan Zhu. 2017. [Assigning personality/identity to a chatting machine for coherent conversation generation](#). *CoRR* abs/1706.02861. <http://arxiv.org/abs/1706.02861>.
- Iulian Vlad Serban, Ryan Lowe, Peter Henderson, Laurent Charlin, and Joelle Pineau. 2015. [A survey of available corpora for building data-driven dialogue systems](#). *CoRR* abs/1512.05742. <http://arxiv.org/abs/1512.05742>.
- Alessandro Sordoni, Michel Galley, Michael Auli, Chris Brockett, Yangfeng Ji, Margaret Mitchell, Jian-Yun Nie, Jianfeng Gao, and Bill Dolan. 2015. [A neural network approach to context-sensitive generation of conversational responses](#). In *Proceedings of the 2015 Conference of the North American Chapter of the Association for Computational Linguistics: Human Language Technologies*. Association for Computational Linguistics, Denver, Colorado, pages 196–205. <http://www.aclweb.org/anthology/N15-1020>.
- P.-H. Su, Y.-B. Wang, T.-H. Yu, and L.-S. Lee. 2013. [A dialogue game framework with personalized training using reinforcement learning for computer-assisted language learning](#). In *2013 IEEE International Conference on Acoustics, Speech and Signal Processing*. pages 8213–8217. <https://doi.org/10.1109/ICASSP.2013.6639266>.
- Sainbayar Sukhbaatar, arthur szlam, Jason Weston, and Rob Fergus. 2015. [End-to-end memory networks](#). In C. Cortes, N. D. Lawrence, D. D. Lee, M. Sugiyama, and R. Garnett, editors, *Advances in Neural Information Processing Systems* 28, Curran Associates, Inc., pages 2440–2448. <http://papers.nips.cc/paper/5846-end-to-end-memory-networks.pdf>.
- Oriol Vinyals and Quoc V. Le. 2015. [A neural conversational model](#). *CoRR* abs/1506.05869. <http://arxiv.org/abs/1506.05869>.
- Min Yang, Zhou Zhao, Wei Zhao, Xiaojun Chen, Jia Zhu, Lianqiang Zhou, and Zigang Cao. 2017. [Personalized response generation via domain adaptation](#). In *Proceedings of the 40th International ACM SIGIR Conference on Research and Development in Information Retrieval*. ACM, New York, NY, USA, SIGIR '17, pages 1021–1024. <https://doi.org/10.1145/3077136.3080706>.
- Weinan Zhang, Ting Liu, Yifa Wang, and Qingfu Zhu. 2017. [Neural personalized response generation as domain adaptation](#). *CoRR* abs/1701.02073. <http://arxiv.org/abs/1701.02073>.

## A Supplemental Material

ACL 2018 also encourages the submission of supplementary material to report preprocessing decisions, model parameters, and other details necessary for the replication of the experiments reported in the paper. Seemingly small preprocessing decisions can sometimes make a large difference in performance, so it is crucial to record such decisions to precisely characterize state-of-the-art methods.

Nonetheless, supplementary material should be supplementary (rather than central) to the paper. **Submissions that misuse the supplementary material may be rejected without review.** Essentially, supplementary material may include explanations or details of proofs or derivations that do not fit into the paper, lists of features or feature templates, sample inputs and outputs for a system, pseudo-code or source code, and data. (Source code and data should be separate uploads, rather than part of the paper).

The paper should not rely on the supplementary material: while the paper may refer to and cite the supplementary material and the supplementary material will be available to the reviewers, they will not be asked to review the supplementary material.

Appendices (*i.e.* supplementary material in the form of proofs, tables, or pseudo-code) should come after the references, as shown here. Use `\appendix` before any appendix section to switch the section numbering over to letters.

## B Multiple Appendices

...can be gotten by using more than one section. We hope you won't need that.